

Just-In-Time compilation of Haskell using PyPy and GHC

Knut Halvor Skrede

December 16, 2011

Abstract

The paper describes a system for using GHC as frontend and PyPy as backend for a Haskell JIT compiler. An intermediate language in JSON format based on GHC's Core language is described. The implemented parts are a serializer written in Haskell and a deserializer written in Python, in addition to some Haskell library functions. The project is meant to serve as a base for further development.

Contents

1	Introduction	4
1.1	Project description	4
1.2	Motivation	4
1.3	System description	4
1.4	Structure of the paper	5
2	Background	5
2.1	PyPy	5
2.1.1	A quick overview	5
2.1.2	RPython	6
2.1.3	Just-in-time compilation	6
2.1.4	Haskell-Python	6
2.2	GHC	6
2.3	Extcore	7
3	External-core	7
3.1	Evaluation of a program	8
4	JSON representation of Core	8
4.1	JSON	8
4.2	JSCore	8
5	Implementation	8
5.1	Tools and versions	8
5.2	Pipeline	9
5.3	Organization	9
5.4	Serializer	10
5.5	Deserializer	11
5.6	Haskell libraries	11
5.7	Evaluation	11
5.8	Tests	12
5.9	Issues	12
6	Example	13
6.1	Example: hello world	13
6.1.1	Converted to Core	14
6.1.2	Converted to JSCore	14
6.1.3	JSCore graph	14
6.1.4	Result	14
7	Future work	15
7.1	Getting programs into Haskell-Python from GHC	15
7.2	Rewrite the deserializer to proper RPython	15
7.3	Implement basic Haskell libraries	15
7.4	Linking functionality	16

7.5 Test suite and benchmarking	16
A Formal definition of External-core	17
B Z-Encoding	20
C Formal definition of JSON	21
D Formal definition of JSCore	22
E Graphs of programs in JSCore	25

List of Figures

1	Project overview	5
2	Code tree: organization	10
3	Pipeline with core2js	11
4	Result output from running tests	12
5	Pipeline without haskell2js	13
6	Example program translated to JSCore	26

List of Tables

2	Grammar for External-core	19
3	Grammar for JSON	21
5	Grammar for JSCore	24

1 Introduction

1.1 Project description

The project is to implement a serializer from Haskell to an intermediate format using GHC, and a deserializer from the intermediate format to an interpreter using PyPy. And then to test the result on some simple Haskell programs. The hope is that this can serve as a base for future development into a full Haskell JIT compiler.

1.2 Motivation

The motivation behind the project is to see if a *strongly-typed non-strict purely-functional* programming language, Haskell, can benefit from just-in-time (JIT) compilation.

Some people seem to think that there is little to gain from using JIT compilation of strongly typed languages since they can be so heavily optimized at compile time. However, a JIT compiler has a lot more information to work with. Speedup in JIT compilers is achieved by feeding values observed at runtime into compilation and exploit them. For example, by locating variables that change very slowly, the JIT compiler can optimize multiple instances of the code, one for each variation of the variable. [2]

PyPy is a project that implements a meta-tracing JIT. The project defines a proper subset of Python called RPython. This language has the characteristics that it is possible to perform type-inference on it. Because of this, interpreters written in RPython can be translated to efficient C code. The idea is that interpreters can be written rapidly in RPython, and the interpreter implemented in RPython will benefit from PyPy's JIT compiler. To optimize the interpreter, the RPython toolchain accepts some compiler hints to determine what parts of the code to trace, and to define the static and dynamic parts of the interpreter "memory". [2]

GHC (Glasgow Haskell Compiler) is the most advanced compiler for the Haskell programming language. The GHC team has defined an external format of Haskell called *External Core*. GHC's front-end translates Haskell 98 (plus some extensions) into Core. This way it is possible to reuse GHC for parsing, desugaring and typechecking, while implementing the back-end separately. [7]

1.3 System description

The system will use the GHC Haskell library to generate the representation of the Core language and a simple Haskell program to serialize this representation into JSON. This entails using GHC for parsing, typechecking, desugaring and

simplification. Simplification is an optimization step in GHC, performing small and simple transformations.

The resulting code will then be deserialized by the PyPy Haskell interpreter and executed. See figure 1 for a description of the implemented and to-be-implemented parts of the project.

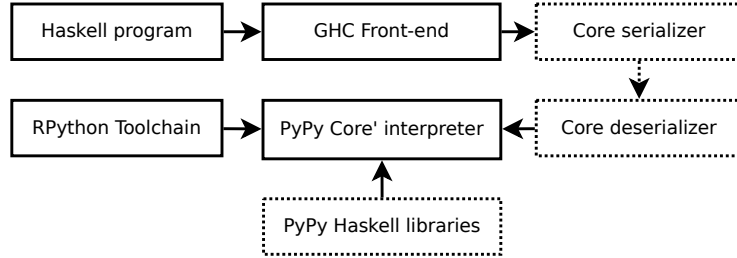


Figure 1: Project overview. Full border: implemented, Dotted border: unimplemented.

1.4 Structure of the paper

Following this introduction is a section discussing the background of the project. The background section contains some information about the tools being used, and why. Section 3 contains a description of the external language representation of the intermediate language used by GHC. Section 4 follows with the description of the intermediate format used in this project, which is a combination of JSON and External-core. Section 5 discusses the implementation of the system, the issues faced during development and some of the remaining issues. In section 6, example programs are passed through the pipeline with intermediate representations presented at each step. Section 7 discusses future work.

2 Background

2.1 PyPy

2.1.1 A quick overview

The PyPy project is basically two things:

1. the RPython toolchain, written in Python, is a set of compiler tools for programs written in RPython.
2. and an implementation of Python using these tools.

In this paper PyPy refers to former.

The basic concept of PyPy is to use a high-level language to allow for rapid development of interpreters for a variety of platforms. By implementing a compiler for RPython, interpreters for other languages can be written in RPython and compiled to any platform supported by the PyPy toolchain. Supported platforms include CLI and JVM. [1]

PyPy uses the meta-programming argument; if a VM (virtual machine) can be written at a level of abstraction high enough, then it should be possible to automatically translate this VM to other lower-level platforms. This is what PyPy does. [6]

2.1.2 RPython

RPython (Restricted Python) is a restricted proper subset of Python that it is possible to perform type inference on. This means that it can be translated to efficient C code, and it enables easy analysis as well as efficient compilation. This also means that RPython code can be run and debugged by Python interpreters, like CPython. [1]

2.1.3 Just-in-time compilation

The JIT compiler is the reason why PyPy is able to compete with other language implementations on speed. Or rather, it's meta-tracing JIT. The JIT is implemented for the RPython compiler, but through a set of compiler hints, it is able to trace the execution of the application interpreted by the RPython program.

2.1.4 Haskell-Python

Haskell-Python is an interpreter for a subset of the Haskell language, called Core'. We call it Core' (Core marked) here because it does not directly correspond to the Core language used by GHC (the glasgow haskell compiler). Haskell-Python is written in RPython, and is compiled into a JIT compiler using the RPython toolchain. Our goal for this project is to extend Haskell-Python to use GHC as a frontend for compilation of Haskell programs.

2.2 GHC

Haskell is a *strongly-typed non-strict purely-functional* programming language, it will not be described in any detail here, since Haskell is not the language we focus on. See [3] for an introduction to Haskell.

Core is an intermediate language used by the Glasgow Haskell Compiler [4], and it is this language we wish to interpret. Core is a desugared version of Haskell, things like pattern matching and list comprehensions are transformed out to simpler constructs. [5]

2.3 Extcore

Extcore is a Haskell package for working with GHC's Core language. Among other things, it implements a parser for External-core, this is the part used from extcore in this project.

3 External-core

The Core language is an intermediate language used by GHC. It is the internal program representation in the compilers simplification phase. External-core is an external representation of Core generated by using a compiler flag. By using External-core, one may implement just parts of a Haskell compiler, using the remaining parts from GHC. For this project, GHC is used to generate External-core. This way, desugaring, type checking, pattern matching and overloading is performed. The remaining task is then to interpret the External-core representation. [7]

The GHC produced External-core uses z-encoding for special characters in names (variables, constructors, ...), for a definition of the z-encoding, see B. This is a good thing, as much of the functionality represented by these names has to be implemented in RPython.

Without using GHC to produce External-core, linking code into GHC would be an optional way of achieving this, which would be a difficult and large task. Or, the GHC API could be used to do the same task more cleanly. [7]

The initial starting point of this project was to use the GHC API, the reasoning was that External-core is not fully parenthesized and more tricky to parse. Thus, generating a fully parenthesized and more machine-readable format would make sense. However, it turned out that this too was a complicated task. As the internal datatypes of GHC representing Core does not match the description of External-core. The choice was then made to use GHC-generated External-core as the base for creating a new intermediate format that could be easily generated using available packages for manipulating External-core.

For a formal definition of External-core, see appendix A.

3.1 Evaluation of a program

A program is evaluated by reducing the expression "main:ZCMain.main" to *weak-head-normal-form* (WHNF), i.e. a primitive value, lambda abstraction, or fully applied data constructor. A heap is used to make sure evaluation is shared. The heap contains two types; a *thunk*, or a *WHNF*. A thunk is an unevaluated expression, also called a *suspension*. A *WHNF* is an evaluated expression, the result of evaluating a *thunk* is a *WHNF* [7]

4 JSON representation of Core

4.1 JSON

JavaScript Object Notation (JSON) is a lightweight data interchange format.

Since a library for manipulating JSON is available for Haskell, this makes it a good choice for the project. In addition, it is easy to parse and the Haskell library contains a pretty-printer, making the result easier to inspect.

For a formal definition of JSON see appendix C

4.2 JSCore

In order to work with JSON and External-core, a format was defined that expresses the Core program in JSON notation. Most of the right hand side of the grammar evaluates to JSON Values. Even though the grammar is changed to support JSON, an effort was made to keep it similar to the original Core grammar for easy referencing. The size of the resulting files was not considered to be an issue.

For a formal definition of JSCore see appendix D

5 Implementation

5.1 Tools and versions

The following tools and package versions was used in the implementation:

- GHC version 7.0.3
- extcore version 1.0.1
- PyPy current head branch (Last tested 09/11/2011)
- Haskell-Python interpreter: Interpreter of Core' written in RPython

- CPython 2.7: Used to test the Core' interpreter without it having to be correct RPython. (pypy-python could also have been used)
- Git was used as the projects version control system.

5.2 Pipeline

The project implements the following pipeline (see figure 3):

1. Serialize Haskell program:
 - (a) Create External-core file from Haskell program using GHC.
 - (b) Create JSCore from External-core using the extcore and JSON packages.
2. Deserialize JSCore:
 - (a) Parse JSCore using the parsing tools available for PyPy
 - (b) Build Core AST from resulting JSON datatype
3. Evaluate program:
 - (a) Evaluation is done by the already implemented PyPy Core' interpreter, Haskell-Python. Additional functionality had to be built on top of this, mostly Haskell library functions.

5.3 Organization

The implementation is organized as represented by figure 2. The main folder (interpreter) contains the main program, and a program for generating dot files, "makegraph.py" (used to create graphs of parsed JSCore files using graphviz). The "haskell" folder contains the PyPy Core' interpreter code, used by the main program for evaluation, and by the parser to generate the abstract-syntax-tree (AST). In addition to this, the subfolder packages implements some simple functionality to be used by the test-programs. Among others, a very simple IO function for printing text to the terminal (putStrLn). These packages are loaded and references to the functions they contain are used during the creation of the AST. See figure 3 for a simple description of the pipeline.

The "core" folder contains a Haskell program for generating JSCore files from External-core files, the JSCore parser, and a simple datastructure representing Haskell modules.

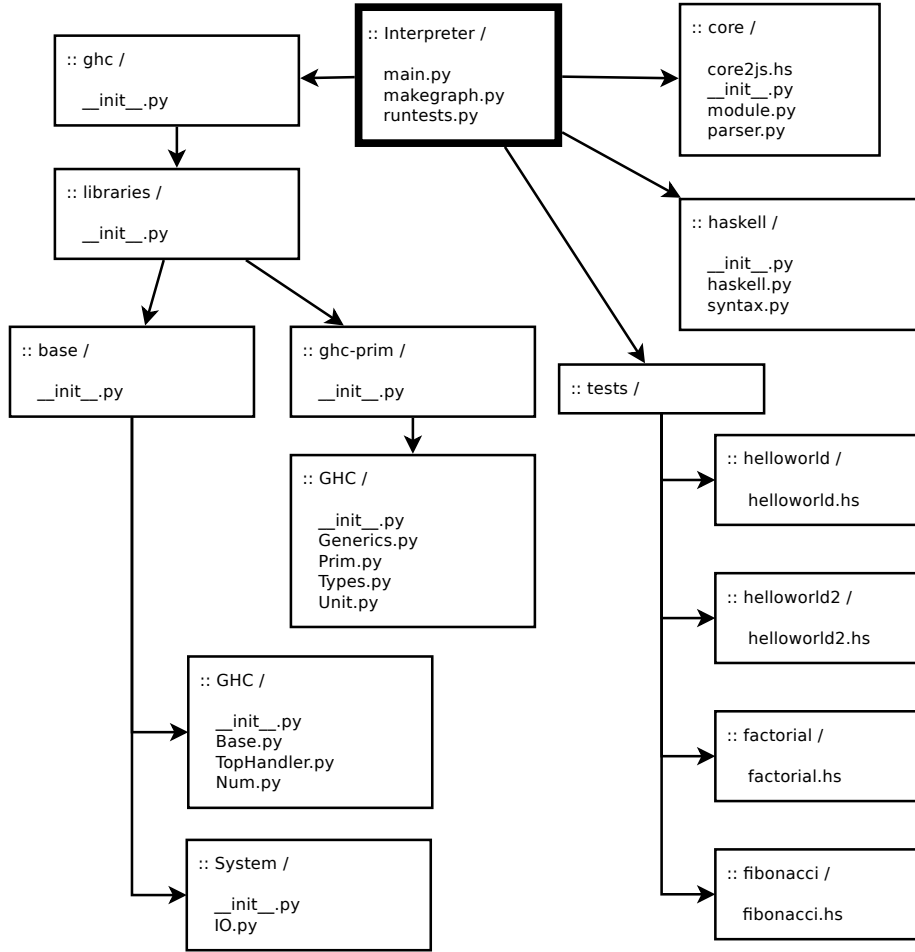


Figure 2: Code tree: The top of the boxes is the folder name, and the rest is the source files. Arrows represent subfolders.

5.4 Serializer

The serializer consists of two parts; GHC generating External-core, and a Haskell program to generate JSCore (core2js).

External-core is easily generated by using a compiler flag:

```
ghc -fext-core {path-to-program}
```

The Haskell program generating JSCore uses the extcore packages. This package implements functionality for working with External-core. The result is a datastructure mapping directly to the External-core format defined in [7]. By

traversing this structure, the program builds up a JSON tree using the Haskell JSON package. The result is a tree of JSON constructs (corresponding to the grammar defined in table 5), this is then pretty-printed and dumped to a file.




Figure 3: Pipeline implementation with core2js using extcore

5.5 Deserializer

The deserializer implements a JSON parser. The resulting datastructure is then traversed, building up an AST using the constructs defined in the Haskell-Python interpreter.

PyPy implements a parser generator, this simply takes a grammar defined as a string, written in extended-backus-naur-form (EBNF), and generates a parser. This parser is then used to create a JSON datastructure, as represented by table 3. The resulting datastructure is then traversed. By checking the contents of the JSON constructs with the actual External-core format, the Core AST is built. External functionality is imported from the Python implementations of the Haskell libraries as it is encountered.

After this is done, we are left with a "module" object, corresponding to the initial Haskell module.

5.6 Haskell libraries

To make some simple test-cases work, some basic Haskell functionality had to be implemented. Some of this functionality was implemented already in the Haskell-Python Core' interpreter. The work done here was mostly to organize the functionality into modules corresponding to Haskell modules. The functionality implemented in these modules does however, not correspond to the Haskell implementations. This is left for future work, as this is a large task.

From figure 6 (representing a "hello world" program in JSCore), the atomic expression "base:SystemziIO.putStrLn" corresponds to the Haskell function "putStrLn", which is located in the Haskell module "System.IO". This is translated into a reference to the function "putStrLn" defined in the python module located in "ghc/libraries/base/System/IO.py". See figure 2.

5.7 Evaluation

In order to evaluate the Haskell programs correctly, the expression "main:ZCMain.main" is reduced to *WHNF*. However, to do this correctly would


require a lot of the functionality used by GHC to be implemented. Specifically, "GHC.TopHandler.runMainIO". In order to implement this function a lot of other functionality would have to be implemented. The function is a wrapper around "main:Main.main", it catches uncaught exceptions and flushes stdout/stderr before exiting.

Currently, a simple hack is implemented; the function "runMainIO" simply return the function "main:Main.main". The "main:Main.main" function is the entry-point to our program. This function is then evaluated when reducing "main:ZCMain.main" to WHNF.

Implementing this is a goal for further development, but currently a simple hack is to only evaluate the expression "main:Main.main". This way, simple programs can be tested by implementing the necessary functionality at a high level, such as the "putStrLn" function which is implemented as a simple "print" Python function.

5.8 Tests

Testing is done by a simple Python script. The test directory has a subdirectory for each test. This directory contains a Haskell program. The program is converted to External-core, then to JSCore, and then executed. If the program returns "0" it passes the test. The results of running these tests are given in figure 4.



RESULTS	
factorial	: Failure
fibonacci	: Failure
helloworld	: Success
helloworld2	: Failure

Figure 4: Result output from running tests (runtests.py)

The tests are supposed to be in the following increasing order of difficulty: helloworld, helloworld2, factorial and fibonacci. All the failures are due to the same problem; unimplemented primitives and library functions.

5.9 Issues

Very few of the initial plans regarding this project was actually realized in the implementation. The GHC API was intended to be used to generate JSCore, but

this failed do to a lack of experience with Haskell and the GHC API (see figure 5 for a description of the initially intended pipeline). It was then discovered that a lot of the necessary code was already written in GHC (the code that generates the External-core files), however, this code was not exported in any way. It was also deeply embedded in the GHC code. After several attempts at creating the JSCore format, and some emails back and forth with the GHC team, these methods where abandoned.




Figure 5: Pipeline with implementation of haskell2js using the GHC API (Not implemented)

The extcore package was then introduced, as it was thought to serve nicely for our purpose. This does however add an extra step to the generation of JSCore, having to generate External-core first. There is no way of working with External-core without parsing it from a file. See figure 3.

The reason for not using extcore from the beginning was that it was thought to not be very well supported, as the External-core format seems to be changing. This also turned out to be the case. However, GHC also changes rapidly, and an implementation using the GHC API may not work for very long either. A version linking to the GHC executable would most likely be the worst choice. As this also changes rapidly.

A large amount of time was spent trying to generate this intermediate format. And the greatest obstacle was that the tools being used was either incompatible, or lacking in documentation. Either way, this problem will have to be revisited and solved.

The method described here worked well very trivial Haskell programs, however, it turned out that the extcore package was not able to parse any nontrivial External-core files generated by the GHC version used. It was thought that this was due to the fact that the extcore package was written for earlier versions of GHC (6.10 and 6.12 and 6.10 on windows). However, these versions turned out to have the same problem when tested.

6 Example

6.1 Example: hello world

The example is a simple "hello world" program, as this was practically the only program that was able to pass through the entire pipeline. Following is the "hello world" program written in Haskell:

```
main = putStrLn "Hello , world!"
```

6.1.1 Converted to Core

After the program has passed through GHC and the External-core file has been generated, the program looks like this:

```
%module main:Main
main:Main.main :: (ghcziprim:GHCziTypes.IO
                  ghcziprim:GHCziUnit.Z0T) =
  base:SystemziIO.putStrLn
  (base:GHCziBase.unpackCStringzh
   ("Hello , world!" :: ghcziprim:GHCziPrim.Addrzh));
main:ZCMain.main :: (ghcziprim:GHCziTypes.IO
                    ghcziprim:GHCziUnit.Z0T) =
  base:GHCziTopHandler.runMainIO @ ghcziprim:GHCziUnit.Z0T
  main:Main.main;
```

6.1.2 Converted to JSCore

In the next step it is parsed and dumped to JSCore:

```
{ "%module": "main:Main", "tdefg": [],
  "vdefg": [ { "qvar": "main:Main.main",
    "ty": { "bty": { "qtycon": "ghcziprim:GHCziTypes.IO" },
    "aty": { "qtycon": "ghcziprim:GHCziUnit.Z0T" } },
    "exp": { "aexp": { "qvar": "base:SystemziIO.putStrLn",
      "args": { "aexp": { "aexp": { "qvar": "base:GHCziBase
        .unpackCStringzh",
        "args": { "aexp": { "lit": "Hello ,
          world!" } } } } },
      "args": { "aexp": { "lit": "Hello ,
        world!" } } } } },
    { "qvar": "main:ZCMain.main",
    "ty": { "bty": { "qtycon": "ghcziprim:GHCziTypes.IO" },
    "aty": { "qtycon": "ghcziprim:GHCziUnit.Z0T" } },
    "exp": { "aexp": { "aexp": { "qvar": "base:GHCziTopHandler.
      runMainIO",
      "args": { "aty": { "qtycon": "ghcziprim:
        GHCziUnit.Z0T" } } },
      "args": { "aexp": { "qvar": "main:Main.main" } } } } } ] }
```

6.1.3 JSCore graph

Using the parsing libraries of PyPy we can generate a nice graph from the result, directly corresponding to the resulting datastructure. See appendix E, figure 6. By simply traversing this datastructure we can generate the AST for the Core' interpreter (Haskell-Python).

6.1.4 Result

The program results as expected, outputting the string "Hello, world!".

7 Future work

7.1 Getting programs into Haskell-Python from GHC

This project focused mainly on this task, which was thought to be simple. However, due to a combination of reasons, this turned out not to be the case. Mainly, inexperience; with Haskell, the GHC API, and functional languages in general. However, a lot of experience was gained during this project, and future development will benefit from this.

The work involved in this will be revisiting the different possibilities to achieve our goal. Some options are:

- Write an External-core parser directly in RPython, and use files generated by GHC.
- Include development of extcore as a part of the project. This is most likely a bad idea, as this seems to be no easier than the alternatives.
- Use the GHC API to generate JSCore, this is also nontrivial.
- Create functionality to be linked into the GHC executable, in order to generate the representation of JSCore. Also nontrivial.
- Simply use the version of GHC matching the version of extcore. This will not be a good idea for further development of the project, but may be a simple solution to get a prototype working quickly.
- Implement a Haskell pipeline in RPython, including parsing, typechecking and desugaring. This will be a very large task.

Either way, understanding of Haskell and Core must be a top priority.

7.2 Rewrite the deserializer to proper RPython

The deserializer is currently not written in proper RPython. Converting the code to RPython will allow it to be compiled to a JIT interpreter by the RPython toolchain. This should not be a very big task, but it requires understanding of the RPython coding style. There are also restrictions on how one may use the PyPy parser tools.

7.3 Implement basic Haskell libraries

Implementing Haskell libraries is necessary to run any Haskell program passed through GHC. One option may be to implement (or automatically generate from GHC code) Haskell primitive types, and to convert the Haskell base libraries to JSCore (or any other format that may be chosen). Compiling GHC with the

`-fext-core` flag should be possible, and should generate External-core files for all library modules.

7.4 Linking functionality

Currently no thought has been given to the linking between multiple modules. This must be implemented for any non-trivial Haskell programs to function.

7.5 Test suite and benchmarking

A framework for testing and benchmarking should be set up. Allowing us to compare speed to other implementations, and progress of development. GHC uses a test framework relying on Python and GNU Make that should be looked into.

A Formal definition of External-core

The following semantics is used to define the Core grammar, as seen in [7]:

[pat]	:	optional
{ pat }	:	zero or more repetitions
{ pat } ⁺	:	one or more repetitions
pat ₁ pat ₂	:	choice

Module

module → %module *mident* { *tdefg* ; } { *vdefg* ; }

Type defn.

<i>tdefg</i>	→	%data <i>qtycon</i> { <i>tbind</i> } = { [<i>cdef</i> { ; <i>cdef</i> }] }	algebraic type
		%newtype <i>qtycon</i> <i>qtycon</i> { <i>tbind</i> } = <i>ty</i>	newtype

Constr. defn.

cdef → *qdcon* { @ *tbind* } { *aty* }⁺

Value defn.

<i>vdefg</i>	→	%rec { <i>vdef</i> { ; <i>vdef</i> } }	recursive
		<i>vdef</i>	non-recursive
<i>vdef</i>	→	<i>qvar</i> :: <i>ty</i> = <i>exp</i>	

Atomic expr.

<i>aexp</i>	→	<i>qvar</i>	variable
		<i>qdcon</i>	data constructor
		<i>lit</i>	literal
		(<i>exp</i>)	nested expr.

Expression

<i>exp</i>	→	<i>aexp</i>	atomic expr.
		<i>aexp</i> { <i>arg</i> } ⁺	application
		\ { <i>binder</i> } -> <i>exp</i>	abstraction
		%let <i>vdefg</i> %in <i>exp</i>	local definition
		%case (<i>aty</i>) <i>exp</i> %of <i>vbind</i> { <i>alt</i> { ; <i>alt</i> } }	case expr.
		%cast <i>exp</i> <i>aty</i>	type coercion
		%note " { <i>char</i> } " <i>exp</i>	expression note
		%external ccal " { <i>char</i> } " <i>aty</i>	external reference
		%dynexternal ccal <i>aty</i>	external reference (dynamic)
		%label " { <i>char</i> } "	external label

Argument

<i>arg</i>	→	@ <i>aty</i>	type argument
		<i>aexp</i>	value argument

Case alt

<i>alt</i>	→	<i>qdcon</i> { @ <i>tbind</i> } { <i>vbind</i> } -> <i>exp</i>	constructor alternative
		<i>lit</i> -> <i>exp</i>	literal alternative
		%_ -> <i>exp</i>	default alternative

Binder

A FORMAL DEFINITION OF EXTERNAL-CORE

<i>binder</i>	→	@ <i>tbind</i>	type binder
		<i>vbind</i>	value binder
Type binder			
<i>tbind</i>	→	<i>tyvar</i>	implicit of kind *
		(<i>tyvar</i> :: <i>kind</i>)	explicitly kinded
Value binder			
<i>vbind</i>	→	(<i>var</i> :: <i>ty</i>)	
Literal			
<i>lit</i>	→	([-] { <i>digit</i> } ⁺ :: <i>ty</i>)	integer
		([-] { <i>digit</i> } ⁺ % { <i>digit</i> } ⁺ :: <i>ty</i>)	rational
		(' <i>char</i> ' :: <i>ty</i>)	character
		(" { <i>char</i> } " :: <i>ty</i>)	string
Character			
<i>char</i>	→	Any ASCII character in range 0x20-0x7E except 0x22, 0x27, 0x5c	
		\x <i>hex hex</i>	ASCII code escape sequence
<i>hex</i>	→	0 ... 9 a ... f	
Atomic type			
<i>aty</i>	→	<i>tyvar</i>	type variable
		<i>qtycon</i>	type constructor
		(<i>ty</i>)	nested type
Basic type			
<i>bty</i>	→	<i>aty</i>	atomic type
		<i>bty aty</i>	type application
		%trans <i>aty aty</i>	transitive coercion
		%sym <i>aty</i>	symmetric coercion
		%unsafe <i>aty aty</i>	unsafe coercion
		%left <i>aty</i>	left coercion
		%right <i>aty</i>	right coercion
		%inst <i>aty aty</i>	instantiation coercion
Type			
<i>ty</i>	→	<i>bty</i>	basic type
		%forall { <i>tbind</i> } ⁺ . <i>ty</i>	type abstraction
		<i>bty</i> -> <i>ty</i>	arrow type construction
Atomic kind			
<i>akind</i>	→	*	lifted kind
		#	unlifted kind
		?	open kind
		<i>bty</i> :=: <i>bty</i>	equality kind
		(<i>kind</i>)	nested kind
Kind			
<i>kind</i>	→	<i>akind</i>	atomic kind
		<i>akind</i> -> <i>kind</i>	arrow kind
Identifier			
<i>mident</i>	→	<i>pname</i> : <i>uname</i>	module
<i>tycon</i>	→	<i>uname</i>	type constr.

<i>qtycon</i>	→	<i>mident . tycon</i>	qualified type constr.
<i>tyvar</i>	→	<i>lname</i>	type variable
<i>dcon</i>	→	<i>uname</i>	data constr.
<i>qdcon</i>	→	<i>mident . dcon</i>	qualified data constr.
<i>var</i>	→	<i>lname</i>	variable
<i>qvar</i>	→	[<i>mident .</i>] <i>var</i>	optionally qualified variable
Name			
<i>lname</i>	→	<i>lower</i> { <i>namechar</i> }	
<i>uname</i>	→	<i>upper</i> { <i>namechar</i> }	
<i>pname</i>	→	{ <i>namechar</i> } ⁺	
<i>namechar</i>	→	<i>lower</i> <i>upper</i> <i>digit</i>	
<i>lower</i>	→	a b ... z _	
<i>upper</i>	→	A B ... Z	
<i>digit</i>	→	0 1 ... 9	

Table 2: Grammar for External-core

B Z-Encoding

The External-core identifiers are z-encoded. Meaning they use "z" and "Z" as prefixes for certain symbols:

Character	Code
Tuples: (##) () (,,)	Z1H Z0T Z3T
Constructors: () [] : Z	ZL ZR ZM ZN ZC ZZ
Variables: z & — ^ \$ = i # . i - ! + , \ / * - % c	zz za zb zc zd ze zg zh zi zl zm zn zp zq zr zs zt zu zv znnnU

C Formal definition of JSON

Object	
<i>object</i>	→ { }
	{ <i>members</i> }
<i>members</i>	→ <i>pair</i>
	<i>pair</i> , <i>members</i>
<i>pair</i>	→ <i>string</i> : <i>value</i>
Array	
<i>array</i>	→ []
	[<i>elements</i>]
<i>elements</i>	→ <i>value</i>
	<i>value</i> , <i>elements</i>
Value	
<i>value</i>	→ <i>string</i>
	<i>number</i>
	<i>object</i>
	<i>array</i>
	true
	false
	null
String	
<i>string</i>	→ ""
	" <i>chars</i> "
<i>chars</i>	→ <i>char</i>
	<i>char</i> <i>chars</i>
<i>char</i>	→ any Unicode character except " or \ or control characters:
	\\
	\/
	\b
	\f
	\n
	\r
	\t
	\u four-hex digits
Number	
<i>number</i>	→ <i>int</i>
	<i>int</i> <i>frac</i>
	<i>int</i> <i>exp</i>
	<i>int</i> <i>frac</i> <i>exp</i>
<i>int</i>	→ <i>digit</i>
	<i>digit</i> 1 – 9 <i>digits</i>
	- <i>digit</i>
	- <i>digit</i> 1 – 9 <i>digits</i>
<i>frac</i>	→ . <i>digits</i>
<i>exp</i>	→ e <i>digits</i>
<i>digits</i>	→ <i>digit</i>
	<i>digit</i> <i>digits</i>
<i>e</i>	→ e
	e+
	e-
	E
	E+
	E-

Table 3: Grammar for JSON

D Formal definition of JSCore

$[pat]$:	Zero or more repetitions of pat surrounded by $[]$ and comma separated (A JSON Array).
$[pat]^+$:	One or more repetitions of pat surrounded by $[]$ and comma separated (A JSON Array).
$\{ pat \}$:	Represents a JSON Object, pat is a JSON <i>members</i> .
$pat_1 \mid pat_2$:	Choice.
$\mid\mid pat \mid\mid$:	Optional

JSCore grammar:

Module			
$module$	\rightarrow	$\{ \%module : mident , \%tdefg : [tdefg] , \%vdefg : [vdefg] \}$	
Type defn.			
$tdefg$	\rightarrow	$\{ \%data : qtycon , \%tbind : [tbind] , \%cdef : [cdef] \}$	algebraic type
	\mid	$\{ \%newtype : qtycon , \%qtycon : qtycon , \%tbind : [tbind] , \%ty : ty \}$	newtype
Constr. defn.			
$cdef$	\rightarrow	$\{ \%qdcon : qdcon , \%tbind : [tbind] , \%aty : [aty]^+ \}$	
Value defn.			
$vdefg$	\rightarrow	$\{ \%rec : [vdef]^+ \}$	recursive
	\mid	$vdef$	non-recursive
$vdef$	\rightarrow	$\{ \%qvar : qvar , \%ty : ty , \%exp : exp \}$	
Atomic expr.			
$aexp$	\rightarrow	$\{ \%qvar : qvar \}$	variable
	\mid	$\{ \%qdcon : qdcon \}$	data constructor
	\mid	$\{ \%lit : lit \}$	literal
	\mid	$\{ \%exp : exp \}$	nested expr.
Expression			
exp	\rightarrow	$aexp$	atomic expr.
	\mid	$\{ \%aexp : aexp , \%args : [arg]^+ \}$	application
	\mid	$\{ \%lambda : [binder] , \%exp : exp \}$	abstraction
	\mid	$\{ \%let : vdefg , \%in : exp \}$	local definition
	\mid	$\{ \%case : aty , \%exp : exp , \%of : vbind , \%alt : [alt]^+ \}$	case expr.
	\mid	$\{ \%cast : exp , \%aty : aty \}$	type coercion
	\mid	$\{ \%note : " \{ char \} " , \%exp : exp \}$	expression note
	\mid	$\{ \%external ccal : " \{ char \} " , \%aty : aty \}$	external reference
	\mid	$\{ \%dynexternal ccal : aty \}$	external reference (dynamic)
	\mid	$\{ \%label : " \{ char \} " \}$	external label
Argument			
arg	\rightarrow	$\{ \%aty : aty \}$	type argument
	\mid	$\{ \%aexp : aexp \}$	value argument
Case alt			
alt	\rightarrow	$\{ \%qdcon : qdcon , \%tbind : [tbind] , \%vbind : [vbind] , \%exp : exp \}$	constructor alternative
	\mid	$\{ \%lit : lit , \%exp : exp \}$	literal alternative
	\mid	$\{ \%_ : exp \}$	default alternative
Binder			

$binder$	\rightarrow	$\{ \text{"tbind"} : tbind \}$	type binder
	$ $	$\{ \text{"vbind"} : vbind \}$	value binder
Type binder			
$tbind$	\rightarrow	$\{ \text{"tyvar"} : tyvar \}$	implicit of kind *
	$ $	$\{ \text{"tyvar"} : tyvar, \text{"kind"} : kind \}$	explicitly kinded
Value binder			
$vbind$	\rightarrow	$\{ \text{"var"} : var, \text{"ty"} : ty \}$	
Literal			
lit	\rightarrow	$jsstring$	string
	$ $	$jsnumber$	number
JSON String			
$jsstring$	\rightarrow	$""$	
	$ $	$\text{" } jschars \text{"}$	
$jschars$	\rightarrow	$jschar$	
	$ $	$jschar jschars$	
$jschar$	\rightarrow	any Unicode character except $"$ or \backslash or control characters: \backslash \backslash $\backslash b$ $\backslash f$ $\backslash n$ $\backslash r$ $\backslash t$ $\backslash u$ four-hex digits	
JSON Number			
$jsnumber$	\rightarrow	$jsint$	
	$ $	$jsint jsfrac$	
	$ $	$jsint jsexp$	
	$ $	$jsint jsfrac jsexp$	
$jsint$	\rightarrow	$jsdigit$	
	$ $	$jsdigit1 - 9 jsdigits$	
	$ $	$- jsdigit$	
	$ $	$- jsdigit1 - 9 jsdigits$	
$jsfrac$	\rightarrow	$. jsdigits$	
$jsexp$	\rightarrow	$jse jsdigits$	
$jsdigits$	\rightarrow	$jsdigit$	
	$ $	$jsdigit jsdigits$	
jse	\rightarrow	e	
	$ $	$e+$	
	$ $	$e-$	
	$ $	E	
	$ $	$E+$	
	$ $	$E-$	
Atomic type			
aty	\rightarrow	$\{ \text{"tyvar"} : tyvar \}$	type variable
	$ $	$\{ \text{"qtycon"} : qtycon \}$	type constructor
	$ $	$\{ \text{"ty"} : ty \}$	nested type
Basic type			
bty	\rightarrow	aty	atomic type
	$ $	$\{ \text{"bty"} : bty, \text{"aty"} : aty \}$	type application
	$ $	$\{ \text{"%trans"} : aty, \text{"aty"} : aty \}$	transitive coercion
	$ $	$\{ \text{"%sym"} : aty \}$	symmetric coercion

D FORMAL DEFINITION OF JSCORE

		{ "%unsafe" : <i>aty</i> , "%aty" : <i>aty</i> }	unsafe coercion
		{ "%left" : <i>aty</i> }	left coercion
		{ "%right" : <i>aty</i> }	right coercion
		{ "%inst" : <i>aty</i> , "%aty" : <i>aty</i> }	instantiation coercion
Type			
	<i>ty</i>	→ <i>btty</i>	basic type
		{ "%forall" : [<i>tbind</i>] ⁺ , "%ty" : <i>ty</i> }	type abstraction
		{ "%bty" <i>btty</i> , "%ty" : <i>ty</i> }	arrow type construction
Atomic kind			
	<i>akind</i>	→ *	lifted kind
		#	unlifted kind
		?	open kind
		{ "%bty" : <i>btty</i> , "%bty" : <i>btty</i> }	equality kind
		{ "%kind" : <i>kind</i> }	nested kind
Kind			
	<i>kind</i>	→ { "%akind" : <i>akind</i> }	atomic kind
		{ "%akind" : <i>akind</i> , "%kind" : <i>kind</i> }	arrow kind
Identifier			
	<i>mident</i>	→ " <i>pname</i> : <i>uname</i> "	module
	<i>tycon</i>	→ " <i>uname</i> "	type constr.
	<i>qtycon</i>	→ " <i>mident</i> . <i>tycon</i> "	qualified type constr.
	<i>tyvar</i>	→ " <i>lname</i> "	type variable
	<i>dcon</i>	→ " <i>uname</i> "	data constr.
	<i>qdcon</i>	→ " <i>mident</i> . <i>dcon</i> "	qualified data constr.
	<i>var</i>	→ " <i>lname</i> "	variable
	<i>qvar</i>	→ " <i>mident</i> . <i>var</i> "	optionally qualified variable
Name			
	<i>lname</i>	→ lower { <i>namechar</i> }	
	<i>uname</i>	→ upper { <i>namechar</i> }	
	<i>pname</i>	→ { <i>namechar</i> } ⁺	
	<i>namechar</i>	→ lower upper <i>digit</i>	
	<i>lower</i>	→ a b ... z _	
	<i>upper</i>	→ A B ... Z	
	<i>digit</i>	→ 0 1 ... 9	

Table 5: Grammar for JSCore

E Graphs of programs in JSCore

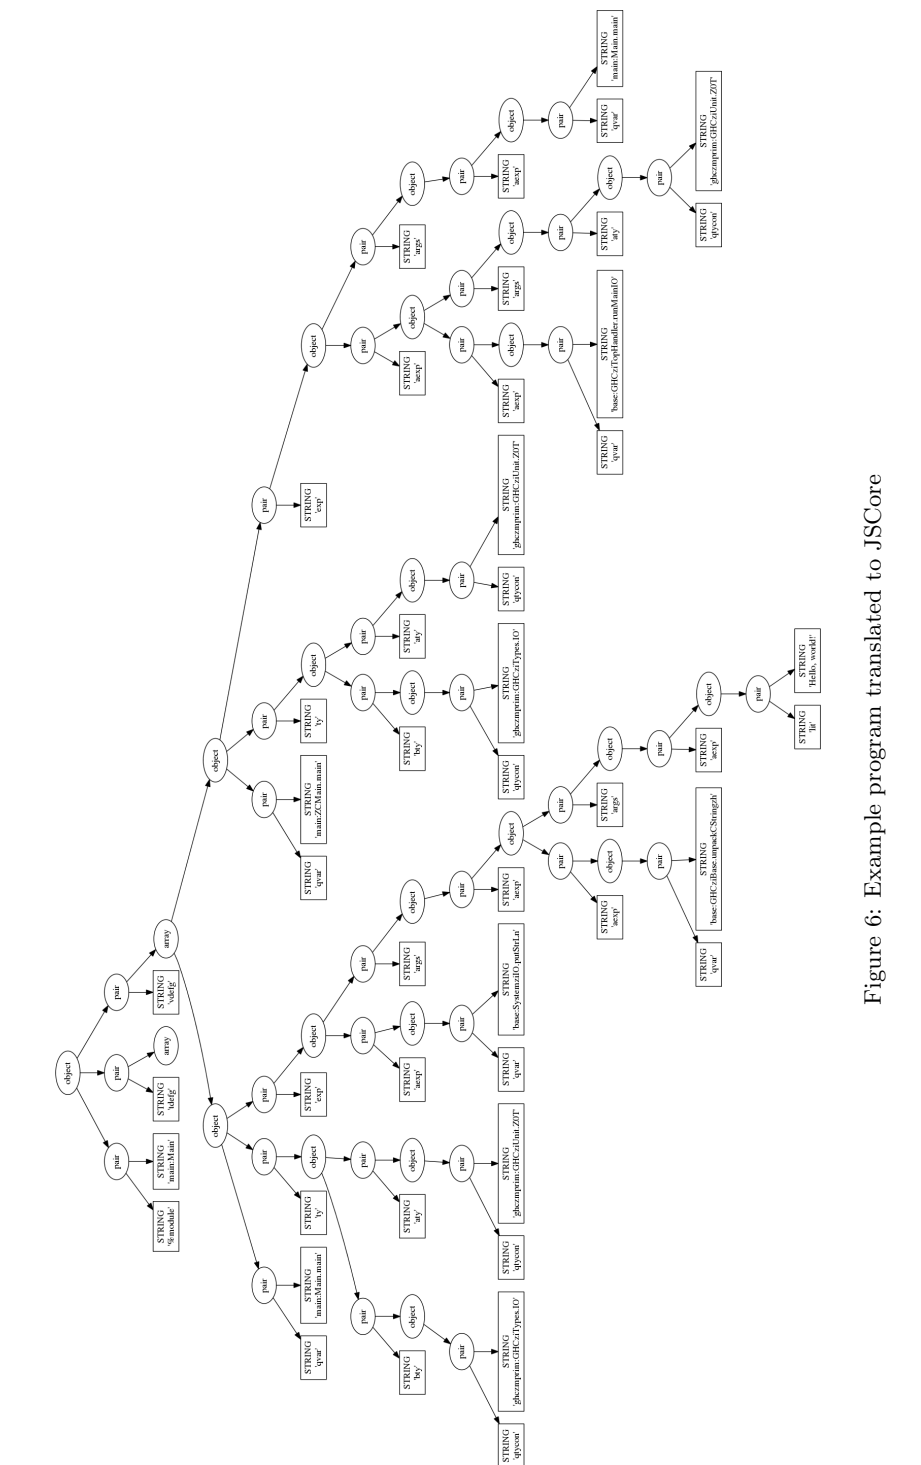


Figure 6: Example program translated to JSCore

References

- [1] D. Ancona, M. Ancona, A. Cuni, and N.D. Matsakis. RPython: a step towards reconciling dynamically and statically typed OO languages. In *Proceedings of the 2007 symposium on Dynamic languages*, pages 53–64. ACM, 2007.
- [2] C.F. Bolz, A. Cuni, M. Fijałkowski, M. Leuschel, S. Pedroni, and A. Rigo. Runtime feedback in a meta-tracing JIT for efficient dynamic languages. In *Proceedings of the 6th Workshop on Implementation, Compilation, Optimization of Object-Oriented Languages, Programs and Systems*, page 9. ACM, 2011.
- [3] P. Hudak, S. Peyton Jones, P. Wadler, B. Boutel, J. Fairbairn, J. Fasel, M.M. Guzmán, K. Hammond, J. Hughes, T. Johnsson, et al. Report on the programming language Haskell: a non-strict, purely functional language version 1.2. *ACM SigPlan notices*, 27(5):1–164, 1992.
- [4] S.L.P. Jones, C. Hall, K. Hammond, J. Cordy, H. Kevin, W. Partain, and P. Wadler. The Glasgow Haskell compiler: a technical overview. 1992.
- [5] S.P. Jones and A. Santos. Compilation by transformation in the Glasgow Haskell Compiler. *Functional Programming, Glasgow*, pages 184–204, 1994.
- [6] A. Rigo and S. Pedroni. PyPy’s approach to virtual machine construction. In *Companion to the 21st ACM SIGPLAN symposium on Object-oriented programming systems, languages, and applications*, pages 944–953. ACM, 2006.
- [7] A. Tolmach and T. Chevalier. the GHC Team. An external representation for the GHC Core language, 2010.